

## AI Interviewer Assignment

*For the Generative AI Engineer Role*

---

### 1. Overview

You are tasked with creating a **full-stack application** that simulates a voice-based AI interviewer. The interviewer should use **Text-to-Speech (TTS)** to ask questions, and **Speech-to-Text (STT)** to transcribe the candidate's responses. An **LLM** (e.g., GPT-4) will provide context-driven follow-up questions and ultimately rate the candidate.

---

### 2. Key Requirements

1. **Real-Time Video Meeting (LiveKit)**
  - Integrate [LiveKit](#) (or an equivalent) for a live video/audio session.
  - The candidate joins the session through a web interface.
  - The AI Interviewer's video feed is optional, but **TTS audio** output is mandatory.
2. **Speech-to-Text (STT) & Text-to-Speech (TTS)**
  - Use **Deepgram** or another STT/TTS provider to:
    1. **Transcribe** candidate responses.
    2. **Speak** the AI Interviewer's questions out loud.
3. **Language Model Integration**
  - Incorporate **GPT-4** (or another LLM) to:
    - Generate interview questions based on the candidate's responses.
    - Provide a **final rating (1–10)** and a **verdict** on candidate suitability.
    - Ensure the **system prompt** can be tailored to different interviewer "personalities."
4. **Local-Only Storage**
  - **No external database** is required.
  - You may store data in memory or a simple local file (e.g., JSON).

- Keep the CV, Job Description, system prompt, transcripts, and results locally.
5. **Frontend & Backend**
- **Frontend** (Admin + Candidate views):
  - **Admin:** Upload or paste the CV, Job Description, and system prompt.
  - **Candidate:** Joins the LiveKit session, speaks, and hears questions via TTS.
  - **Backend:**
  - Orchestrates calls to GPT-4, STT/TTS, and LiveKit.
  - Stores interview data and final results.
6. **End-to-End Flow**
- **Admin** configures the interview by providing the candidate's details (CV, JD, system prompt).
  - **Candidate** joins the interview URL.
  - AI Interviewer greets them with **TTS**, listens via **STT**, and uses **GPT-4** to decide subsequent questions.
  - At the end, the AI Interviewer generates a **rating** and **verdict**, stored locally.
- 

### 3. Deliverables

1. **Source Code**
  - A Git repository (or ZIP) with your entire solution.
  - Both **frontend** (admin and candidate views) and **backend** (orchestration logic) should be included.
2. **README**
  - **Setup Instructions:** How to install dependencies, run the application locally, and set environment variables (e.g., OPENAI\_API\_KEY, DEEPGRAM\_API\_KEY).
  - **Usage:** Step-by-step guide to simulate an interview (admin side + candidate side).
  - **Notes:** Any additional comments, known issues, or limitations.
3. **Brief Demo/Explanation**

- If possible, a short video or written walkthrough demonstrating a sample interview session (optional but appreciated).
- 

#### 4. Evaluation Criteria

1. **Code Quality**
    - Readable, well-structured code.
    - Proper separation of concerns and use of clear naming conventions.
  2. **Implementation Accuracy**
    - Correct usage of STT/TTS and GPT-4.
    - Minimal dependencies beyond what is necessary.
  3. **User Experience**
    - Candidates can easily join the LiveKit room, speak, and hear TTS responses.
    - Admin interface is straightforward for uploading CV, JD, prompt, and viewing results.
  4. **Completeness**
    - The system should produce a **final rating** and short **verdict**.
    - All data (CV, JD, transcripts, results) should be stored and retrievable locally.
  5. **Documentation**
    - A clear README that allows others to run and test your solution without confusion.
- 

#### 5. Submission

- **Submission:**
- Send us a link to your Git repository or a ZIP archive.
- Include the README file at the root level.
- **Questions:** If you have any clarifications, email [akash@growthut.in](mailto:akash@growthut.in)